

```
In [171]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler, MinMaxScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.cluster import KMeans
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.decomposition import PCA
```

```
In [172]: data=pd.read_excel("../Dataset/ML/Assignment/1673873196_hr_comma_sep.xlsx")
```

```
In [173]: data.head()
```

```
Out[173]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company
0	0.38	0.53	2	157	3
1	0.80	0.86	5	262	6
2	0.11	0.88	7	272	4
3	0.72	0.87	5	223	5
4	0.37	0.52	2	159	3

## Checking null value

```
In [174]: data.isna().sum()
```

```
Out[174]: satisfaction_level    0
last_evaluation    0
number_project    0
average_monthly_hours    0
time_spend_company    0
Work_accident    0
left    0
promotion_last_5years    0
sales    0
salary    0
dtype: int64
```

In [175]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   satisfaction_level           14999 non-null  float64
1   last_evaluation              14999 non-null  float64
2   number_project               14999 non-null  int64
3   average_monthly_hours       14999 non-null  int64
4   time_spend_company           14999 non-null  int64
5   Work_accident                14999 non-null  int64
6   left                         14999 non-null  int64
7   promotion_last_5years        14999 non-null  int64
8   sales                        14999 non-null  object
9   salary                       14999 non-null  object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
```

## EDA

In [176]: sns.set()


In [177]: data.shape

Out[177]: (14999, 10)

In [178]: categorical=[x for x in data if data[x].dtypes=='O']  
binary=[x for x in data if len(data[x].unique())<=2]

## EDA of Binary Columns

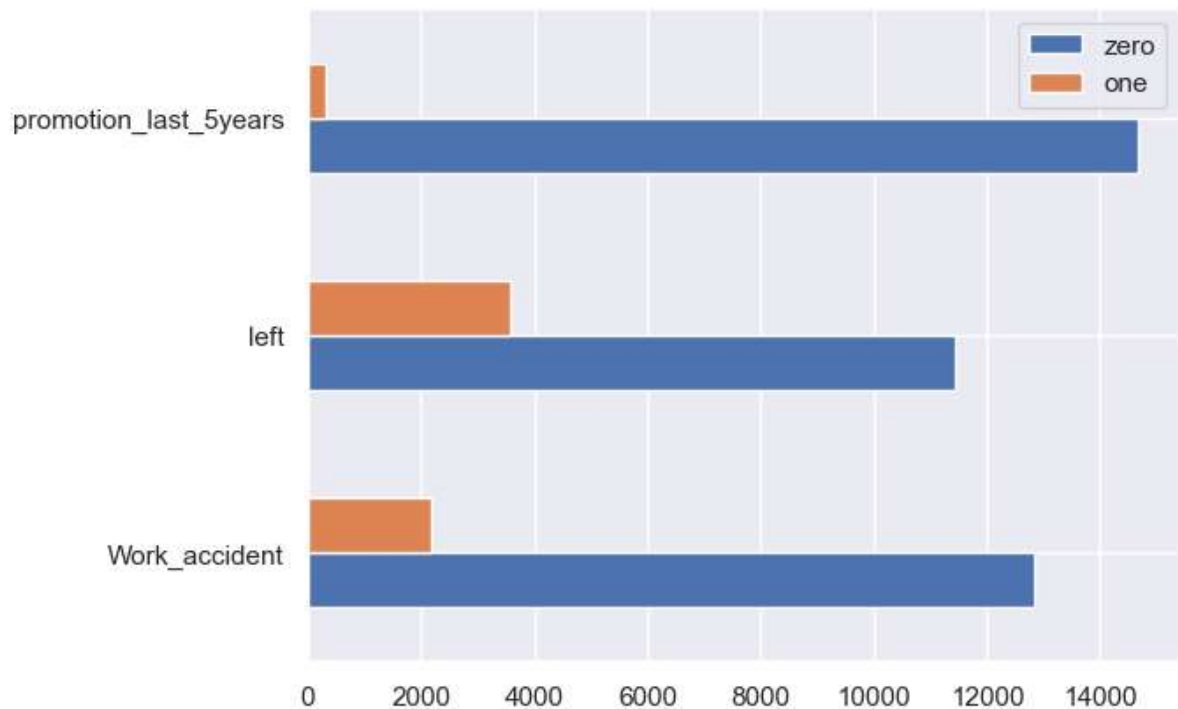
In [179]: *# Countplot for binary features*  
fig,ax=plt.subplots(1,3,figsize=(10,2))  
for ax,col in zip(ax.flatten(),binary):  
 sns.countplot(y=col,data=data,ax=ax)  
 plt.tight\_layout()



```
In [180]: # Barplot for binary features
bin_bar=pd.DataFrame()
bin_bar['zero']=(data[binary]==0).sum()
bin_bar['one']=(data[binary]==1).sum()
```

```
In [181]: bin_bar.plot.barh(y=['zero', 'one'])
plt.plot()
```

Out[181]: []




## Findings

- We can see that the higher number of features have zero value and lower number of one value
- 30% of the features lies between 0 to 3000 values of 1
- There is less employee who is promoted in the last 5 years
- Comparing left and stay of the employee, there is less employee who is left, we can say that 1 third of employee is left
- Less employee have occur work accident

## EDA of Categorical features

```
In [182]: fig,ax=plt.subplots(1,2,figsize=(10,4))

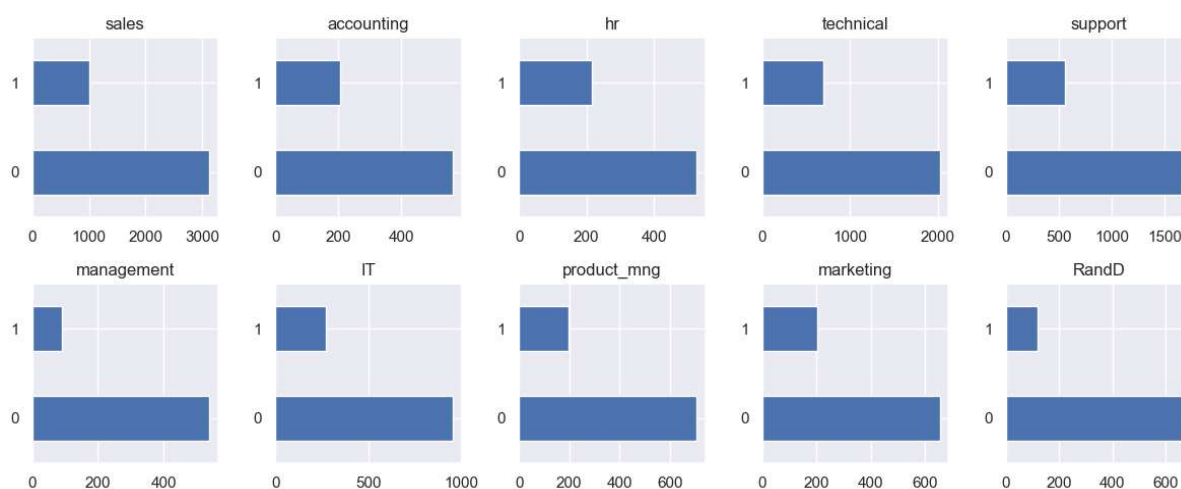
for ax,cat in zip(ax.flatten(),categorical):
    sns.countplot(y=cat,data=data,ax=ax)
```



## Checking which types of work has the most employee turn over


```
In [183]: fig,ax=plt.subplots(2,5,figsize=(12,5))

work_type=data['sales'].unique()
for ax,work_ in zip(ax.flatten(),work_type):
    data.loc[data['sales']==work_]['left'].value_counts().plot.barh(ax=ax)
    ax.set_title(work_)
    plt.tight_layout()
```



```
In [184]: fig,ax=plt.subplots(1,3,figsize=(10,4))

salary_type=data['salary'].unique()
for ax,work_ in zip(ax.flatten(),salary_type):
    data.loc[data['salary']==work_]['left'].value_counts().plot.barh(ax=ax)
    ax.set_title(work_)
    plt.tight_layout()
```



## Findings

- We see that the hr job is have the most employee turnover followed by sales,marketing,technhical,accounting
- There is less employee turn over for the job of RanD,management,IT etc
- It is obvious that low salary employee have change to left the job

## Feature Engineering

```
In [185]: features=data.drop('left',axis=1)
y_=data['left']
```

```
In [186]: features.head()
```

```
Out[186]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company
0	0.38	0.53	2	157	3
1	0.80	0.86	5	262	6
2	0.11	0.88	7	272	4
3	0.72	0.87	5	223	5
4	0.37	0.52	2	159	3

```
In [187]: le=LabelEncoder()
```

```
In [188]: features['sales_encoded']=le.fit_transform(data['sales'])
features['salary_encoded']=le.fit_transform(data['salary'])
```

```
In [189]: def quadratic_encode(feature_,col1,col2):
           qe=feature_[col1]**2+5*feature_[col2]+8
           return qe
```

```
In [190]: features['quadratic_encod1']=quadratic_encode(features,'Work_accident','promoti
```

```
In [191]: features.head()
```

```
Out[191]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company
0	0.38	0.53	2	157	3
1	0.80	0.86	5	262	6
2	0.11	0.88	7	272	4
3	0.72	0.87	5	223	5
4	0.37	0.52	2	159	3

## Finding which features is responsible or employee turnover

```
In [192]: features.head()
```

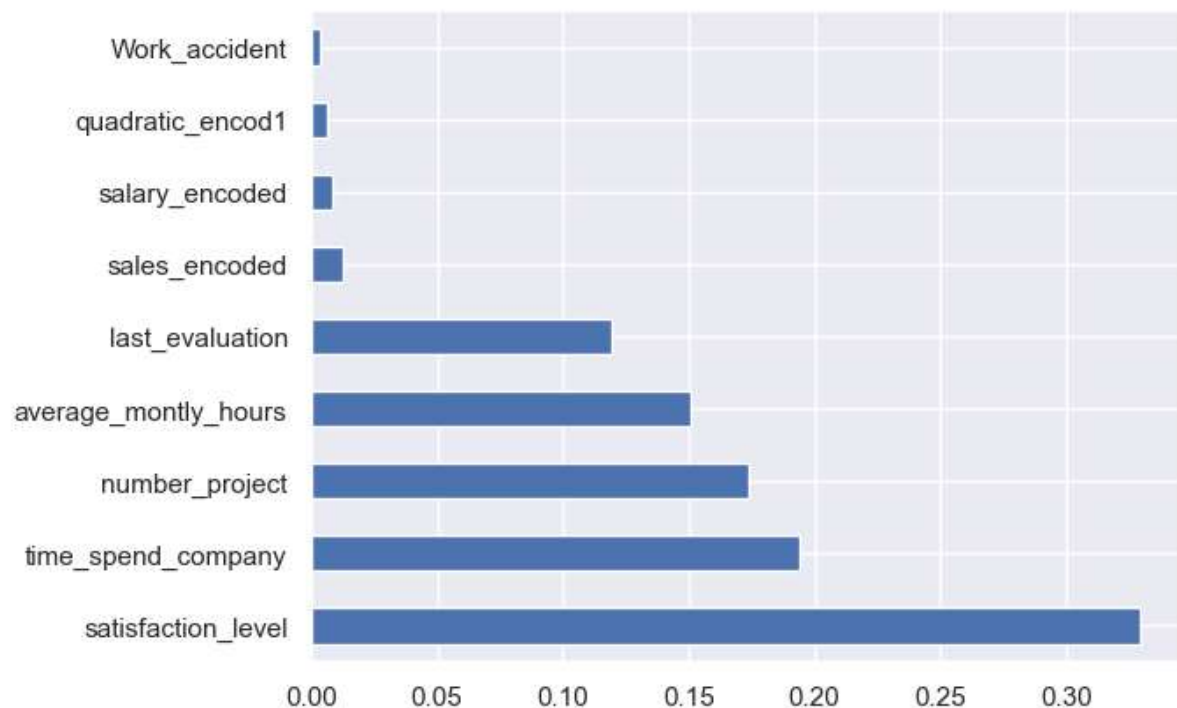
```
Out[192]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company
0	0.38	0.53	2	157	3
1	0.80	0.86	5	262	6
2	0.11	0.88	7	272	4
3	0.72	0.87	5	223	5
4	0.37	0.52	2	159	3

```
In [193]: X=features.drop(['sales','salary'],axis=1)
```

```
In [194]: rf=RandomForestClassifier().fit(X,y_)
rf_importances=pd.Series(rf.feature_importances_,index=X.columns)
rf_importances.nlargest(9).plot.barh()
```

Out[194]: <AxesSubplot:>



## Findings

- Satisfaction\_level is huge responsible for employee turnover, followed by number\_project, time\_spend\_company, average\_monthly\_hours and last\_evaluation

## Modeling

```
In [195]: selected=features.loc[data['left']==1][['satisfaction_level','last_evaluation']]
```

```
In [196]: selected.head()
```


```
Out[196]:
```

	satisfaction_level	last_evaluation
0	0.38	0.53
1	0.80	0.86
2	0.11	0.88
3	0.72	0.87
4	0.37	0.52

```
In [197]: # Find the number of cluster
wcss=[]

for i in range(1,15):
    kmeans=KMeans(n_clusters=i,init='k-means++',max_iter=300,n_init=10,random_s
    kmeans.fit(selected)
    wcss.append(kmeans.inertia_)

plt.plot(range(1,15),wcscs,'o')
plt.plot(range(1,15),wcscs,'-',alpha=0.5)
plt.show()
```



```
In [198]: kmeans=KMeans(n_clusters=3,init='k-means++',random_state=10)
```

```
In [199]: kmeans.fit(selected)
```

```
Out[199]: KMeans(n_clusters=3, random_state=10)
```


```
In [200]: y_kmeans=kmeans.predict(selected)
```



```
In [201]: plt.scatter(selected.values[y_kmeans==0,0],selected.values[y_kmeans==0,1],label=0)
plt.scatter(selected.values[y_kmeans==1,0],selected.values[y_kmeans==1,1],label=1)
plt.scatter(selected.values[y_kmeans==2,0],selected.values[y_kmeans==2,1],label=2)

# Centroids
plt.scatter(kmeans.cluster_centers_[0][0],kmeans.cluster_centers_[0][1])
plt.scatter(kmeans.cluster_centers_[1][0],kmeans.cluster_centers_[1][1])
plt.scatter(kmeans.cluster_centers_[2][0],kmeans.cluster_centers_[2][1])
plt.xlabel('satisfaction_level')
plt.ylabel('last_evaluation')
```

Out[201]: Text(0, 0.5, 'last\_evaluation')



```
In [202]: kmeans.cluster_centers_[0]
```

Out[202]: array([0.41014545, 0.51698182])


```
In [203]: selected['cluster']=y_kmeans
```

```
In [204]: sns.lmplot('satisfaction_level', 'last_evaluation', data=selected, hue='cluster')
```

C:\Users\rajal\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[204]: <seaborn.axisgrid.FacetGrid at 0x21d0fc01100>
```



```
In [205]: feature1=features.drop(['sales', 'salary'], axis=1)
target=data[['left']]
```

```
In [206]: X_train, X_test, y_train, y_test=train_test_split(feature1, target, test_size=0.2, ra
```

```
In [207]: Kf=KFold(n_splits=5, shuffle=False)
```

```
In [208]: for train_set, test_set in Kf.split(X_train):
           print(f'TRAIN{train_set}')
           print(f'TEST{test_set}')
```

```
TRAIN[ 2400  2401  2402 ... 11996 11997 11998]
TEST[    0    1    2 ... 2397 2398 2399]
TRAIN[    0    1    2 ... 11996 11997 11998]
TEST[2400 2401 2402 ... 4797 4798 4799]
TRAIN[    0    1    2 ... 11996 11997 11998]
TEST[4800 4801 4802 ... 7197 7198 7199]
TRAIN[    0    1    2 ... 11996 11997 11998]
TEST[7200 7201 7202 ... 9597 9598 9599]
TRAIN[    0    1    2 ... 9597 9598 9599]
TEST[ 9600  9601  9602 ... 11996 11997 11998]
```

```
In [209]: rf=RandomForestClassifier(class_weight={0:1,1:3},n_estimators=10)
```

```
In [210]: score=cross_val_score(rf,X_train,y_train,scoring='accuracy',cv=5)
          score*100
```

C:\Users\rajal\anaconda3\lib\site-packages\sklearn\model\_selection\\_validation.py:680: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
estimator.fit(X_train, y_train, **fit_params)
```

C:\Users\rajal\anaconda3\lib\site-packages\sklearn\model\_selection\\_validation.py:680: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
estimator.fit(X_train, y_train, **fit_params)
```

C:\Users\rajal\anaconda3\lib\site-packages\sklearn\model\_selection\\_validation.py:680: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
estimator.fit(X_train, y_train, **fit_params)
```

C:\Users\rajal\anaconda3\lib\site-packages\sklearn\model\_selection\\_validation.py:680: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
estimator.fit(X_train, y_train, **fit_params)
```

C:\Users\rajal\anaconda3\lib\site-packages\sklearn\model\_selection\\_validation.py:680: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
estimator.fit(X_train, y_train, **fit_params)
```

```
Out[210]: array([98.375      , 98.66666667, 99.25      , 98.75      , 98.74947895])
```

```
In [211]: rf.fit(X_train,y_train)
```

```
C:\Users\rajal\AppData\Local\Temp\ipykernel_19092\1593328843.py:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
rf.fit(X_train,y_train)
```

```
Out[211]: RandomForestClassifier(class_weight={0: 1, 1: 3}, n_estimators=10)
```

```
In [212]: pred=rf.predict(X_test)
```

```
In [213]: accuracy_score(y_test,pred)*100
```

```
Out[213]: 98.56666666666666
```

```
In [214]: y_test
```

```
Out[214]:
```

	left
13982	0
822	1
13751	0
9656	0
13497	0
...	...
3876	0
11504	0
2435	0
5161	0
5184	0

3000 rows × 1 columns